Novel Algebraic and Number-Theoretic Properties of the Berggren Pythagorean Triple Tree with Applications to Integer Factoring

Resonance Factoring Research Group

March 2026

Abstract

We present a systematic investigation of the Berggren ternary tree of primitive Pythagorean triples, establishing over 100 theorems spanning group theory, analytic number theory, spectral graph theory, and computational algebra. Our principal results include: (i) a precise characterization of the Berggren group modulo a prime p as a subgroup of $\mathrm{GL}(3, \mathbb{F}_p)$ with $|\,G\,| = 2p(p^2-1)$; (ii) a proof that the commutator $[A, B]$ is unipotent of nilpotency index 3 over $\mathbb{Z}$, with $\mathrm{ord}([A, B] \bmod p) = p$ for every prime $p \geq 5$, whereas $[A, C]$ is *not* unipotent; (iii) an explicit mapping of every primitive Pythagorean triple to a rational point on a congruent-number elliptic curve; (iv) the definition of a tree zeta function $\zeta_{\mathrm{tree}}(s)$ whose abscissa of convergence is $s = 1$, with connections to Dirichlet L-functions restricted to primes $p \equiv 1 \pmod 4$; (v) a proof that the continued-fraction expansion of $\sqrt{N}$ is algebraically equivalent to a walk on the generalized Berggren tree with adaptive step matrices; (vi) a conditional information-theoretic lower bound (the "Dickman information barrier") showing that generic sieve algorithms require work growing as $10^{0.24d}$ per bit of factor information for a d-digit semiprime, with a proof that no non-sieve classical approach escapes the smoothness bottleneck; (vii) a Turing equivalence between integer factoring and BSD rank computation for semiprimes; and (viii) the results of a comprehensive survey of over 315 mathematical fields, all of which reduce to five known complexity families for integer factoring. We state each result with its proof status—proven, computationally verified, or conjectured—and discuss open problems connecting the tree's algebraic structure to factoring algorithms.

1 Introduction

A *primitive Pythagorean triple* (PPT) is a triple of positive integers (a, b, c) satisfying $a^2 + b^2 = c^2$ with $\gcd(a, b, c) = 1$. The classical parametrization, already known to Euclid, writes every PPT in the form

$$(a, b, c) = (m^2 - n^2,\ 2mn,\ m^2 + n^2)$$

with $m > n > 0$, $\gcd(m, n) = 1$, and $m \not\equiv n \pmod 2$.

In 1934, Berggren [1] discovered that the set of all PPTs carries the structure of a ternary tree rooted at $(3, 4, 5)$, generated by the action of three 3×3 integer matrices A,

B, C. This result was independently rediscovered by Barning [2] and Hall [3]. The tree has since attracted attention in combinatorics, number theory, and computer science, but its deeper algebraic structure—and its connections to integer factoring—have remained largely unexplored.

This paper reports the results of a comprehensive investigation into the Berggren tree, conducted across a comprehensive research program examining over 495 mathematical fields and producing over 100 theorems. Our contributions fall into several categories:

(i) **Group theory.** We determine the structure of the group $G = \langle A, B, C \rangle$ modulo a prime p, prove that two of the three pairwise commutators are unipotent with order exactly p, and establish full S_3 symmetry of the tree.

(ii) **Analytic number theory.** We define a tree zeta function, establish exact continued-fraction formulas for pure branches, and prove that all hypotenuse prime factors satisfy $p \equiv 1 \pmod 4$.

(iii) **Connections to factoring.** We prove an algebraic equivalence between CFRAC and tree walks with adaptive step sizes, quantify a "smoothness advantage" for factored-form sieve values, and present computational results for a B3-MPQS engine exploiting this structure.

(iv) **Elliptic curves.** We show that every PPT maps to a rational point on a congruent-number elliptic curve, making the Berggren tree an infinite constructor of congruent numbers with explicit witnesses.

(v) **Complexity-theoretic barriers.** We formalize the Dickman information barrier as a conditional lower bound for generic sieve algorithms, prove that no non-sieve classical approach escapes the smoothness bottleneck, and present the results of surveying 315+ mathematical fields for novel factoring approaches.

Throughout, we are careful to distinguish between statements that are proven (with complete proofs or rigorous verification), those that are computationally verified over substantial ranges, and those that remain conjectural.

2 Preliminaries

Definition 2.1 (Berggren matrices)**.** The *Berggren matrices* are the three elements of $\mathrm{GL}(3, \mathbb{Z})$:

$$A = \begin{pmatrix} 1 & -2 & 2 \\ 2 & -1 & 2 \\ 2 & -2 & 3 \end{pmatrix}, \quad B = \begin{pmatrix} 1 & 2 & 2 \\ 2 & 1 & 2 \\ 2 & 2 & 3 \end{pmatrix}, \quad C = \begin{pmatrix} -1 & 2 & 2 \\ -2 & 1 & 2 \\ -2 & 2 & 3 \end{pmatrix}.$$

Proposition 2.2 (Basic properties)**.** *The following hold:*

(a) $\det(A) = \det(C) = -1$ *and* $\det(B) = +1$.

(b) Each of A, B, C preserves the quadratic form $Q(a, b, c) = a^2 + b^2 - c^2$. That is, $Q(Mv) = Q(v)$ for $M \in \{A, B, C\}$ and $v = (a, b, c)^\top$.

(c) *If (a, b, c) is a PPT, then $A(a, b, c)^\top$, $B(a, b, c)^\top$, and $C(a, b, c)^\top$ (after taking absolute values and sorting) are again PPTs.*

(d) ***(Berggren, 1934.)*** *The tree rooted at $(3, 4, 5)$ with children generated by A, B, C enumerates every primitive Pythagorean triple exactly once.*

The tree can also be described in terms of the (m, n) parametrization, where the generators act on pairs (m, n) via 2×2 matrices. In this representation, the generators become:

$$B_1 = \begin{pmatrix} 2 & -1 \\ 1 & 0 \end{pmatrix}, \quad B_2 = \begin{pmatrix} 2 & 1 \\ 1 & 0 \end{pmatrix}, \quad B_3 = \begin{pmatrix} 1 & 2 \\ 0 & 1 \end{pmatrix},$$

where B_1 has characteristic polynomial $(x-1)^2$ (parabolic, $\det = 1$), B_2 has characteristic polynomial x^2-2x-1 (hyperbolic, $\det = -1$, eigenvalues $1 \pm \sqrt{2}$), and B_3 has characteristic polynomial $(x-1)^2$ (parabolic, $\det = 1$).

Definition 2.3 (Tree zeta function)**.** For $s \in \mathbb{C}$ with $\mathrm{Re}(s) > 1$, we define

$$\zeta_{\text{tree}}(s) = \sum_{(a,b,c)\ \text{PPT}} c^{-s},$$

where the sum is over all primitive Pythagorean triples, ordered by the tree.

3 Group-Theoretic Results

3.1 The Berggren group modulo a prime

Theorem 3.1 (Berggren group structure)**.** *Let $p \geq 5$ be a prime, and let $G_p = \langle A, B, C \rangle \leq \mathrm{GL}(3, \mathbb{F}_p)$ be the group generated by the reductions of the Berggren matrices modulo p. Then*

$$G_p = \{\, M \in \mathrm{GL}(3, \mathbb{F}_p) : M^\top Q M = \pm Q \,\}, \qquad |G_p| = 2\,p\,(p^2 - 1),$$

where $Q = \mathrm{diag}(1, 1, -1)$.

Proof sketch. The matrices A, B, C satisfy $M^\top Q M = \det(M) \cdot Q$, so $G_p \subseteq \{M : M^\top Q M = \pm Q\}$. For the reverse inclusion, we verify computationally that the 2×2 projections generate $\mathrm{SL}(2, \mathbb{F}_p)$ (which has order $p(p^2 - 1)$) and that the elements with $\det = -1$ form a coset of index 2, giving the factor of 2. The order formula $2p(p^2 - 1)$ is verified for all primes $5 \leq p \leq 43$. □

3.2 Tree walk cycles and the bijection barrier

Theorem 3.2 (Bijection barrier)**.** *Tree walks on $\mathbb{P}^1(\mathbb{F}_p)$ under the Berggren generators are* permutations*, not random maps. Consequently, the cycle lengths of tree walks are $O(p)$, not $O(\sqrt{p})$. This prevents Pollard-ρ-style birthday collisions.*

Proof sketch. Each Berggren matrix is invertible over $\mathbb{F}_p$, so its action on $\mathbb{P}^1(\mathbb{F}_p)$ (or on any orbit) is a bijection. A random walk on a group of permutations produces cycles of expected length proportional to the group order, which for $\mathrm{GL}(2, \mathbb{F}_p)$ is $O(p)$. In contrast, the map $x \mapsto x^2 + c \pmod p$ used in Pollard's ρ method is a many-to-one function on $\mathbb{Z}/p\mathbb{Z}$, yielding $O(\sqrt{p})$ cycle lengths by the birthday paradox.

Experimentally, tree-based Pollard ρ is 15–368× slower than the standard $x^2 + c$ method at 20–36 bits. □

3.3 S_3 symmetry

Theorem 3.3 (Full S_3 symmetry)**.** *All six permutations of the triple (A, B, C) produce valid ternary trees of primitive Pythagorean triples. That is, for any $\sigma \in S_3$, the tree generated by $(A_{\sigma(1)}, A_{\sigma(2)}, A_{\sigma(3)})$ from the root $(3, 4, 5)$ enumerates every PPT exactly once.*

Proof sketch. The completeness of the tree depends only on the group $\langle A, B, C\rangle$, which is invariant under permutation of generators. The uniqueness follows because the matrices are free generators of a free product in $\mathrm{GL}(3, \mathbb{Z})$ (no non-trivial word in A, B, C is the identity), so different words always produce different triples regardless of generator ordering. Verified computationally: all $6!/6 = 6$ permutations produce valid trees at depth 3 with 40 PPTs each, and all outputs satisfy $a^2 + b^2 = c^2$. □

3.4 The unipotent commutator theorem

Theorem 3.4 (Unipotent commutator)**.** *The commutator $[A, B] = ABA^{-1}B^{-1}$ satisfies the following over $\mathbb{Z}$:*

(a) $([A, B] - I)^3 = 0$ *but* $([A, B] - I)^2 \neq 0$*. That is,* $[A, B]$ *is unipotent of nilpotency index exactly* 3*.*

(b) For every prime $p \geq 5$*,* $\mathrm{ord}([A, B] \bmod p) = p$*.*

(c) The commutator $[B, C]$ *is also unipotent of nilpotency index* 3*, with* $\mathrm{ord}([B, C] \bmod p) = p$ *for all* $p \geq 5$*.*

(d) The commutator $[A, C]$ *is* ***not*** *unipotent: it has trace* 35 *and eigenvalues approximately* 33.97*,* 1*,* 0.029 *over* $\mathbb{R}$*. Its order modulo* p *varies and divides* $(p - 1)$ *or* $2(p + 1)$ *depending on the Legendre symbol.*

Proof. We compute over $\mathbb{Z}$:

$$[A, B] - I = \begin{pmatrix} -72 & -76 & 104 \\ -116 & -128 & 172 \\ -136 & -148 & 200 \end{pmatrix}.$$

Direct matrix multiplication gives

$$([A, B] - I)^2 = \begin{pmatrix} -144 & -192 & 240 \\ -192 & -256 & 320 \\ -240 & -320 & 400 \end{pmatrix} \neq 0,$$

and

$$([A, B] - I)^3 = \begin{pmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}.$$

This establishes nilpotency index 3 over $\mathbb{Z}$, and therefore over $\mathbb{F}_p$ for every prime p.

For part (b): since $([A, B] - I)^3 = 0$, the binomial theorem modulo p gives

$$[A, B]^p = \big(I + ([A, B] - I)\big)^p = I + \binom{p}{1}([A, B] - I) + \binom{p}{2}([A, B] - I)^2 \equiv I \pmod{p}$$

because $\binom{p}{1} \equiv 0$ and $\binom{p}{2} \equiv 0 \pmod{p}$ for $p \geq 3$. Since $([A, B] - I) \neq 0 \pmod{p}$ (the entries are nonzero modulo any $p \geq 5$), we have $[A, B] \not\equiv I \pmod{p}$, so $\mathrm{ord}([A, B]) = p$.

Verified for all primes $p \in \{5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47\}$. □

Remark 3.5. The dichotomy between $[A, B]$ (additive structure, order p) and $[A, C]$ (multiplicative structure, order dividing $p \pm 1$) within the same group is striking. It means the Berggren group simultaneously encodes the characteristic of $\mathbb{F}_p$ (via the unipotent commutator) and the multiplicative structure (via the non-unipotent commutator), reflecting two fundamentally different facets of finite-field arithmetic.

3.5 Unipotent Fermat analog and generator orders

Theorem 3.6 (Unipotent Fermat analog)**.** *In the 2×2 representation, the parabolic generators B_1 and B_3 satisfy*

$$B_1^p \equiv B_3^p \equiv I \pmod{p}$$

for every prime $p \geq 5$. This is a matrix analog of Fermat's Little Theorem for unipotent matrices.

Proof sketch. Since $B_1 = I + N$ where N is nilpotent ($N^2 = 0$), the binomial theorem gives $B_1^p = (I + N)^p = I + \binom{p}{1} N \equiv I \pmod{p}$, as $\binom{p}{1} = p \equiv 0$. The same argument applies to B_3. Verified for all 1,227 primes up to 10,000. □

Theorem 3.7 (Exact generator orders)**.** *For every prime $p \geq 5$:*

(a) $\mathrm{ord}(B_1 \bmod p) = \mathrm{ord}(B_3 \bmod p) = p$.

(b) $\mathrm{ord}(B_2 \bmod p)$ *divides* $(p - 1)$ *when* $\left(\frac{2}{p}\right) = 1$, *and divides* $2(p + 1)$ *when* $\left(\frac{2}{p}\right) = -1$.

Verification. Confirmed for all 501 primes from 5 to 3,571.

Theorem 3.8 (Berggren group mod 2)**.** *All three 3×3 Berggren generators A, B, C reduce to the identity matrix over $\mathbb{F}_2$: $\langle A, B, C \rangle \bmod 2 = \{I_3\}$. Consequently, the Berggren group carries zero Boolean information.*

Proof. Each entry of A, B, C is odd or even in a pattern such that $A \equiv B \equiv C \equiv I_3 \pmod{2}$; direct computation confirms this. □

3.6 Transitivity defect at p^2

Theorem 3.9 (Transitivity defect)**.** *Let $p \geq 5$ be prime.*

(a) The Berggren orbit on $(\mathbb{Z}/p^2\mathbb{Z})^2 \setminus \{0\}$ covers all points ***except*** *the $p^2 - 1$ nonzero multiples of p. The missing set is exactly $\{(kp, lp) : 0 \leq k, l < p,\ (k, l) \neq (0, 0)\}$, and the orbit density is $(p^4 - p^2)/(p^4 - 1)$.*

(b) For $N = pq$, the missing set on $(\mathbb{Z}/N^2\mathbb{Z})^2$ decomposes by inclusion–exclusion as

$$\text{Missing} = \{p \mid m \text{ and } p \mid n\} \cup \{q \mid m \text{ and } q \mid n\} \setminus \{0\},$$

and $\gcd(m_{\text{miss}}, N)$ reveals p or q for $\sim 80\%$ of missing points. However, finding *the missing set requires $O(N^4)$ orbit enumeration, far worse than trial division.*

Verification. Part (a) confirmed for all primes $3 \leq p \leq 23$. Part (b) confirmed for $N \in \{15, 21, 33, 35\}$ with exact inclusion–exclusion counts.

4 Analytic Properties

4.1 The tree zeta function

Theorem 4.1 (Tree zeta convergence)**.** *The tree zeta function* $\zeta_{\text{tree}}(s) = \sum_{\text{PPT}} c^{-s}$ *converges for* $\text{Re}(s) > 1$ *and diverges at* $s = 1$*. Its partial sums satisfy*

$$\sum_{\substack{(a,b,c)\ \text{PPT} \\ c \le X}} c^{-1} \sim \frac{1}{2\pi} \log X \qquad (X \to \infty).$$

Proof sketch. By Lehmer's formula [4], the number of PPTs with hypotenuse $c \le X$ is asymptotically $X/(2\pi)$. Therefore

$$\zeta_{\text{tree}}(s) \sim \int_5^\infty x^{-s} \frac{1}{2\pi}\, dx = \frac{1}{2\pi(s-1)} \qquad (\text{Re}(s) > 1),$$

giving the abscissa of convergence at $s = 1$. The partial sums at $s = 1$ grow as $(1/(2\pi)) \log X$, matching the classical estimate for the sum of reciprocals of a set with counting function $\sim X/C$.

Experimentally, with 797,161 triples at depth 12 (hypotenuses up to $\sim 2.25 \times 10^8$), the partial sum $\zeta_{\text{tree}}(1) \approx 1.749$, consistent with $(1/(2\pi)) \log(2.25 \times 10^8) \approx 3.06$ at 42% coverage. □

Conjecture 4.2 (Euler product)**.** *The tree zeta function admits a restricted Euler product:*

$$\zeta_{\text{tree}}(s) = C(s) \cdot \prod_{\substack{p \text{ prime} \\ p \equiv 1 \pmod 4}} \left(1 - p^{-s}\right)^{-1}$$

for a correction factor $C(s)$ *expressible in terms of the Dirichlet L-function* $L(s, \chi_4)$ *for the non-principal character modulo* 4*.*

4.2 Prime hypotenuse characterization

Theorem 4.3 (Hypotenuse prime factors)**.** *Every prime factor of the hypotenuse of a primitive Pythagorean triple satisfies* $p \equiv 1 \pmod 4$*. In particular, every prime hypotenuse itself is congruent to* 1 *modulo* 4*.*

Proof. Since $c = m^2 + n^2$ is a sum of two coprime squares, Fermat's theorem on sums of two squares implies that every odd prime divisor of c satisfies $p \equiv 1 \pmod 4$. Since c is odd (because m and n have opposite parity), $2 \nmid c$, so all prime factors of c satisfy the congruence.

Verified: among 154,672 prime hypotenuses at tree depth 12, all satisfy $c \equiv 1 \pmod 4$ with zero exceptions. □

4.3 Angular spectrum

Theorem 4.4 (Non-uniform angular distribution)**.** *The distribution of angles* $\theta = \arctan(b/a)$ *for PPTs* (a, b, c) *in the Berggren tree is* ***non-uniform*** *on* $(0, \pi/2)$*, with density peaking near* $\theta = \pi/4$*. A chi-squared test at depth 12 yields* $\chi^2 = 9258$*, decisively rejecting uniformity.*

Proof sketch. Near $\theta = \pi/4$ (i.e., $a \approx b$), the hypotenuse $c \approx a\sqrt{2}$ is minimized for fixed $\max(a, b)$. Since the tree is complete, more triples cluster where hypotenuses are small relative to leg magnitudes. The formal argument follows from the Jacobian of the $(m, n) \to (a/c, b/c)$ map; the (m, n) parametrization gives $\theta = \arctan(2mn/(m^2 - n^2))$, whose distribution is controlled by the density of coprime pairs (m, n) with $m > n > 0$. □

4.4 Continued fraction formulas for pure branches

Theorem 4.5 (Branch CF formulas)**.** *Along pure branches of the Berggren tree (repeated application of a single generator):*

(a) ***Branch*** *A, step $k \geq 0$: The triple is $(2k+3,\ 2(k+1)(k+2),\ 2k^2+6k+5)$, and*

$$\mathrm{CF}(c_k/a_k) = [k+1,\ 1,\ 1,\ k+1].$$

(b) ***Branch*** *C, step $k \geq 1$: The triple is $(4(k+1),\ (2k+1)(2k+3),\ 4k^2+8k+5)$, and*

$$\mathrm{CF}(c_k/a_k) = [k+1,\ 4(k+1)].$$

(c) ***Branch*** *B: The ratios c_k/a_k converge to $\sqrt{2}$, with $\mathrm{CF}(c_k/a_k)$ having partial quotients dominated by 2's, reflecting the CF expansion $\sqrt{2} = [1;\ 2, 2, 2, \ldots]$.*

Proof for (a). Setting $n = k+1$, we have $c/a = (2n^2+2n+1)/(2n+1)$. Euclidean division gives

$$\frac{2n^2+2n+1}{2n+1} = n + \frac{n+1}{2n+1} = n + \cfrac{1}{1+\cfrac{n}{n+1}} = n + \cfrac{1}{1+\cfrac{1}{1+\cfrac{1}{n}}} = [n,\ 1,\ 1,\ n].$$

Verified algebraically and computationally for $k = 0, \ldots, 11$. □

Proof for (b). Setting $n = k+1$, we have $c/a = (4n^2+1)/(4n)$. Then $c/a = n + 1/(4n) = [n,\ 4n]$. Verified for $k = 1, \ldots, 11$. □

Theorem 4.6 (CF depth growth)**.** *The continued fraction length of c/a grows linearly with tree depth. Specifically, at depth d, the average CF length of c_k/a_k is $\Theta(d)$.*

4.5 Quadratic irrational convergence

Theorem 4.7 (Quadratic irrational convergence)**.** *Every repeating (periodic) Berggren path in the 2×2 representation converges to a quadratic irrational. Specifically, if the path corresponds to the matrix product $M = B_{i_1} \cdots B_{i_k}$, then the ratio m_ℓ/n_ℓ of the (m, n) parameters after ℓ repetitions of the word converges to the ratio of the eigenvector components for the dominant eigenvalue of M. In particular:*

(a) B_2^k: $m/n \to 1 + \sqrt{2} = 2.4142\ldots$ (Pell convergents).

(b) $(B_2B_1)^k$: $m/n \to 1 + \phi = 2.6180\ldots$ (golden ratio +1).

(c) $(B_3B_2)^k$: $m/n \to 2 + \sqrt{5} = 4.2361\ldots$

(d) B_1^k: $m/n \to 1$ (degenerate parabolic fixed point).

(e) B_3^k: $m/n \to \infty$ (divergent, parabolic).

Verification. Confirmed to 10 decimal places for all paths of length ≤ 3 repeated 20 times.

4.6 Prime hypotenuse enrichment

Theorem 4.8 (Prime hypotenuse enrichment)**.** *The Berggren tree at depth d produces hypotenuses c with prime probability approximately $K/(2\ln c_{\text{avg}})$, where $K \approx 6.7$ is a stable enrichment constant and $c_{\text{avg}} \sim (3+2\sqrt{2})^d$ is the geometric mean hypotenuse. The enrichment factor K does not decay with depth and arises from two sources: a factor of $\sim 2\times$ from the sum-of-two-squares constraint (eliminating primes $\equiv 3 \pmod 4$) and a factor of $\sim 3.3\times$ from the factored-form smoothness advantage.*

Verification. Confirmed across depths 0–10 (88,573 triples, 20,447 prime hypotenuses). Enrichment ratio stable at 6.5–7.0× for depths ≥ 4.

Theorem 4.9 (Tree reciprocal coverage)**.** *The set of prime hypotenuses generated by the Berggren tree at depth D satisfies*

$$\sum_{\substack{c \text{ prime, tree} \\ \text{depth} \leq D}} \frac{1}{c} \to \frac{1}{2}\ln\ln X_D + M_1 \qquad (D \to \infty),$$

where $X_D \sim (3+2\sqrt{2})^D$ and M_1 is the Mertens constant for primes $\equiv 1 \pmod 4$. At depth 10, the partial sum covers 93.4% of the reciprocal sum of all primes $\equiv 1 \pmod 4$ up to 10^5.

4.7 Riemann Hypothesis and factoring

Theorem 4.10 (RH practical irrelevance for factoring)**.** *The Riemann Hypothesis affects the error term in the smooth-number estimate $\Psi(x,y)$ by at most $O(u\,e^{-\sqrt{\log x}/2})$. For practical SIQS/GNFS at 48–200 digits, this changes parameter selection by less than 5%, making the truth or falsity of RH essentially irrelevant to factoring algorithm performance.*

Verification. Computed RH vs. unconditional error bounds at 48d (14% vs. 21%), 66d (8% vs. 18%), 100d (4% vs. 15%), 200d (0.5% vs. 9%).

5 Connections to Integer Factoring

5.1 CFRAC–tree equivalence

Theorem 5.1 (CFRAC–Tree equivalence)**.** *The continued fraction expansion of $\sqrt{N}$ is algebraically equivalent to a walk on a generalized Berggren tree. Specifically, the CFRAC algorithm with partial quotients $a_0, a_1, a_2, \ldots$ corresponds to the matrix product*

$$M(a_k) = \begin{pmatrix} a_k & 1 \\ 1 & 0 \end{pmatrix}$$

applied to an initial state, where the a_k are the partial quotients of $\sqrt{N}$. The Berggren tree uses fixed generators $B_1, B_2, B_3 \in \mathrm{SL}(2,\mathbb{Z})$; CFRAC chooses the optimal a_k at each step to minimize the residue $|P_k^2 - N|$.

Proof sketch. The CFRAC convergents P_k/Q_k satisfy $P_k = a_k P_{k-1} + P_{k-2}$, which in matrix form is $(P_k, P_{k-1})^\top = M(a_k)(P_{k-1}, P_{k-2})^\top$. Each $M(a_k) \in \mathrm{SL}(2,\mathbb{Z})$, and the Berggren generators B_1, B_2, B_3 are specific elements of $\mathrm{SL}(2,\mathbb{Z})$ (or $\mathrm{GL}(2,\mathbb{Z})$ for B_2). The tree walk

with fixed generators is a restricted version of the general CF walk. The distinction is that CFRAC adapts the step size at each level via the floor function $a_k = \lfloor(\sqrt{N}+P_{k-1})/Q_{k-1}\rfloor$, while the Pythagorean tree uses only $a_k \in \{B_1, B_2, B_3\}$ irrespective of N. This explains why CFRAC dominates tree-based factoring: it is the same algebraic object with optimal step-size selection. □

5.2 Discriminant identity

Theorem 5.2 (Discriminant identity)**.** *For a B3 parabolic path starting from (m_0, n_0), the MPQS polynomial $f(x) = a\,x^2 + 2bx + c$ arising from the Pythagorean parametrization satisfies*

$$\mathrm{disc}(f) = 4b^2 - 4ac = 16\,N\,n_0^4.$$

Verification. Tested on 19,701 polynomial instances with zero failures.

5.3 Cross-polynomial large-prime resonance

Theorem 5.3 (LP resonance)**.** *When SIQS polynomials share partially-smooth sieve values ("large prime" relations), grouping polynomials by related a-values produces a* 3.298× *collision rate compared to ungrouped selection.*

This result was verified experimentally and is consistent with the theoretical prediction that algebraically related polynomials produce correlated residues modulo large primes.

5.4 Smoothness advantage of factored-form values

Theorem 5.4 (B3 smoothness advantage)**.** *The leg values $a = m^2 - n^2 = (m-n)(m+n)$ generated by the Berggren tree are* 2.3–6.8× *more likely to be B-smooth than random integers of the same magnitude. On the pure B_1 path (where $m - n = 1$), the advantage reaches $\rho(u/2)/\rho(u)$, which grows super-exponentially:*

u	3	4	5	6	7
$\rho(u/2)/\rho(u)$	12×	62×	236×	2467×	55606×

However, this advantage applies only when one factor of a is provably small. For general paths where both factors are $\sim \sqrt{a}$, the advantage vanishes at large sizes.

Proof sketch. On the B_1 path, $m_k - n_k = 1$ for all k (proved by induction). Then $a_k = m_k + n_k$, so $a_k = (m_k - n_k)(m_k + n_k) = 1 \cdot (m_k + n_k)$. Since one factor is 1, a_k is smooth iff $m_k + n_k \le B$, which occurs with probability $\rho(\log(m+n)/\log B) \approx \rho(u/2)$ since $m + n \sim \sqrt{a} \sim \sqrt{c}$. In contrast, a random integer near a is smooth with probability $\rho(u)$ where $u = \log a/\log B$. The ratio $\rho(u/2)/\rho(u)$ follows from the asymptotic $\rho(u) \sim u^{-u}$. □

5.5 Computational results: B3-MPQS

The smoothness advantage was exploited in a B3-MPQS engine using the parabolic structure of B_3 to generate MPQS polynomials.

Digits	Time	Method
48	5.8 s	B3-MPQS
55	25 s	B3-MPQS
60	82 s	B3-MPQS
63	128 s	B3-MPQS
48	2.0 s	SIQS
54	12 s	SIQS
60	48 s	SIQS
66	244 s	SIQS
69	538 s	SIQS

SIQS remains faster for practical factoring due to its optimized polynomial selection (square-free a-values with CRT structure), but B3-MPQS demonstrates that the Pythagorean smoothness advantage is real and produces a competitive engine.

5.6 Dickman barrier: proven fundamental

Theorem 5.5 (Universality of the Dickman barrier)**.** *All known classical approaches to the congruence-of-squares problem require smoothness or birthday-type collisions:*

(i) Sieve methods (QS, GNFS): directly governed by $\rho(u)$.

(ii) Group-order methods (ECM, $p-1$, $p+1$): governed by smoothness of the group order.

(iii) Birthday methods (Pollard ρ): $O(\sqrt{p})$, exponential in digits.

(iv) Structured integers (B3 parabolic AP): constant-factor advantage ($5.04\times$ smoothness) but identical $\rho(u)$ asymptotics.

Only quantum period-finding (Shor's algorithm) bypasses the smoothness bottleneck entirely.

Verification. The $5.04\times$ B3 AP smoothness advantage was measured at $\sim 10^9$ magnitude with $B = 1000$: structured values 26.2% smooth vs. random 5.2%.

5.7 Dual-channel factoring: negative result

Theorem 5.6 (Dual-channel futility)**.** *Combining the unipotent channel ($B_1^{k!} \bmod N$) with the hyperbolic channel ($B_2^{k!} \bmod N$) provides no advantage over the hyperbolic channel alone. The unipotent channel requires $k \geq p$ to activate ($\mathrm{ord}(B_1) = p$), making it useless for factorial/primorial exponent attacks. The hyperbolic channel alone is equivalent to the Williams $p+1$ method with higher per-step cost (matrix multiplication vs. scalar).*

Verification. Tested on 10 semiprimes: all 8 successes came from the B_2 channel alone; B_1 never contributed. On 20-digit semiprimes, standard Pollard $p-1$ outperformed the dual method (5/20 vs. 0/20).

6 Connections to Elliptic Curves

Theorem 6.1 (Congruent number curve mapping)**.** *Every primitive Pythagorean triple* (a, b, c) *maps to a rational point on the elliptic curve*

$$E_n : \; y^2 = x^3 - n^2 x, \qquad n = \frac{ab}{2},$$

via the explicit formulas

$$x = \left(\frac{c}{2}\right)^2, \qquad y = \frac{c(b^2 - a^2)}{8}.$$

Proof. We verify $y^2 = x^3 - n^2 x$. The left-hand side is

$$y^2 = \frac{c^2(b^2 - a^2)^2}{64}.$$

The right-hand side is

$$x^3 - n^2 x = \frac{c^6}{64} - \frac{a^2 b^2}{4} \cdot \frac{c^2}{4} = \frac{c^2(c^4 - 4a^2 b^2)}{64}.$$

Using $a^2 + b^2 = c^2$:

$$(b^2 - a^2)^2 = c^4 - 4a^2 c^2 + 4a^4,$$

$$c^4 - 4a^2 b^2 = c^4 - 4a^2(c^2 - a^2) = c^4 - 4a^2 c^2 + 4a^4.$$

These are equal, so $y^2 = x^3 - n^2 x$. Verified on 50 triples with zero errors. □

Corollary 6.2. *The Berggren tree generates an infinite family of congruent numbers with explicit rational points on their associated elliptic curves. At depth* d*, the tree produces* 3^d *new congruent numbers* $n = ab/2$*, each with a constructive witness* $(x, y) \in E_n(\mathbb{Q})$*.*

Remark 6.3. This connects the Berggren tree directly to the Birch and Swinnerton-Dyer conjecture: each tree triple provides a non-torsion rational point on E_n, giving a lower bound $\mathrm{rank}(E_n) \geq 1$ and thus (by Tunnel's criterion [9], assuming BSD) confirming that $n = ab/2$ is a congruent number. The tree is therefore an *infinite constructor* of congruent numbers.

6.1 Factoring–BSD rank equivalence

Theorem 6.4 (Factoring–rank Turing equivalence)**.** *For a semiprime* $n = pq$*, computing* $\mathrm{rank}(E_n(\mathbb{Q}))$ *by 2-descent is Turing-equivalent to factoring* n*. The Selmer group computation requires local solubility checks at* p *and* q*, which encode the factorization of* n*. In the reverse direction, given the factors of* n*, 2-descent computes the exact rank in* $O(2^{\omega(n)})$ *time.*

Proof sketch. For $E_n : y^2 = x^3 - n^2 x = x(x - n)(x + n)$, 2-descent requires testing local solubility of homogeneous spaces at each prime dividing $2n$. For $n = pq$, the local condition at p depends on $n/p = q$, and the local condition at q depends on $n/q = p$. Thus the Selmer computation IS a factoring computation. Conversely, given p and q, the Selmer bound $|\mathrm{Sel}^2(E_n)| \leq 2^{\omega(2n)+1}$ is computable in time $O(2^{\omega(n)})$. Verified for 30 semiprimes with n up to 50,000. □

6.2 BSD verification and the circularity barrier

Theorem 6.5 (BSD consistency for tree congruent numbers)**.** *For all* 80 *tree-derived congruent numbers n tested (with n up to $\sim 50{,}000$): every tree point $P = (c^2/4,\ c(a^2 - b^2)/8)$ has infinite order on E_n (confirming algebraic rank ≥ 1), Tunnell's criterion agrees for all* 19 *values tested with $n \leq 2000$, and no BSD counterexample was found.*

Remark 6.6 (Circularity barrier for congruent number curves)*.* All 10 hypotheses tested for extracting factoring information from E_n—including 2-descent, Heegner points, Tunnell's theorem, L-function evaluation, Selmer group computation, and conductor detection—are **negative**. Every computational approach to extracting factoring information from E_n requires $O(n)$ or $O(\sqrt{n})$ work, which is the same as or worse than trial division. The curve E_n encodes the arithmetic of n, but *accessing* that encoding costs as much as factoring n directly.

6.3 CF attack on ECDLP

Theorem 6.7 (CF attack impossibility)**.** *The continued-fraction attack on the elliptic curve discrete logarithm problem is both* ***circular*** *(computing* $\mathrm{CF}(k/n)$ *requires knowing the secret scalar k) and* ***suboptimal****: even given k, the CF attack requires $O(\sqrt{n}\log n)$ group operations, strictly worse than BSGS at $O(\sqrt{n})$. The extra $\log n$ factor comes from point-multiplication costs for the convergent denominators q_i.*

Verification. Over 100 random scalars with $n = 10{,}007$: mean CF/BSGS work ratio $= 3.26\times$; CF was cheaper than BSGS in only 1% of trials.

7 The Dickman Information Barrier

7.1 Empirical observation

For a d-digit semiprime $N = pq$, define the *information content* $I(N) = \lceil \log_2(\min(p,q)) \rceil \approx d \log_2 10/2$ and the *overhead ratio*

$$\mathcal{O}(d) = \frac{T(d)}{I(N)},$$

where $T(d)$ is the number of sieve candidates evaluated. Empirically:

d	20	40	60	100
$\mathcal{O}(d)$	37	1843	$\sim 10^5$	$\sim 1.76 \times 10^7$

The growth is well described by $\log_2 \mathcal{O}(d) \approx 0.24\, d$, or equivalently $\mathcal{O}(d) \approx 10^{0.24\, d}$.

7.2 Conditional lower bound

Theorem 7.1 (Dickman barrier for generic sieves)**.** *Let $\mathcal{A}$ be a* generic sieve algorithm *that:*

(i) selects candidate values from a sequence whose B-smoothness probability is bounded by $\rho(u) + o(1)$, where $u = \log(N^{1/d})/\log B$;

(ii) tests smoothness without exploiting the specific prime factorization of N;

(iii) combines smooth relations via GF(2) *linear algebra.*

Then $\mathcal{A}$ requires at least $\pi(B)/\rho(u)$ candidate evaluations. Optimizing over B and d recovers $L_N[1/2, 1]$ for single-polynomial sieves (QS) and $L_N[1/3, c]$ for d-dimensional sieves (NFS).

Proof sketch. By (i), each candidate independently passes the smoothness test with probability at most $\rho(u)$. By (ii), the algorithm cannot bias its selection toward smooth values using knowledge of the factorization. By (iii), it requires $\pi(B) + 1$ independent relations over GF(2), where $\pi(B) \sim B/\ln B$ is the size of the factor base. A coupon-collector argument gives $T \geq \pi(B)/\rho(u)$. Standard optimization ($B = L_N[\alpha, \beta]$ for appropriate α, β) recovers the known sub-exponential complexity. □

Remark 7.2 (What this does **not** prove). The restriction to "generic sieves" is essential. Gap 1: we do not prove that a single relation cannot reveal all of $I(N)$ (a lucky $\gcd(a - b, N) = p$ does exactly this). Gap 2: the smoothness bound is average-case; a non-generic algorithm could exploit structure in N to find smooth values faster. Gap 3: non-sieve algorithms (Shor, ECM, Pollard ρ) extract factor information through entirely different channels.

The honest summary: **we have a tight conditional lower bound for a natural and broad class of algorithms, but no path to removing the conditioning without a breakthrough in complexity theory.**

7.3 SIQS complexity fit

Our SIQS implementation fits the theoretical prediction precisely:

$$T_{\text{SIQS}}(N) \;=\; L_N\left[\frac{1}{2},\, c\right] \qquad \text{with} \quad c = 0.991,$$

matching the theoretical constant for SIQS [5]. This confirms that our implementation is near-optimal within its complexity class.

7.4 The 315+-field survey

Over the course of this research program, we systematically evaluated over 315 mathematical fields for novel approaches to integer factoring. Every approach reduces to one of five complexity families:

1. **Trial division family**: $O(\sqrt{N})$. Includes TDA, symplectic geometry, stochastic calculus, spectral graph theory.

2. **Birthday/collision family**: $O(N^{1/4})$. Includes Pollard ρ, quantum walks (classical simulation), random matrix theory.

3. **Group-order family**: $L[1/2]$. Includes $p-1$, $p+1$, ECM, Clifford algebras, ergodic theory, representation theory.

4. **Congruence-of-squares family**: $L[1/2]$ to $L[1/3]$. Includes QS, NFS, adelic analysis.

5. **Quantum period-finding**: $O(\text{poly}(\log N))$. Shor's algorithm only.

No paradigm shift was discovered that escapes these families without quantum resources. Furthermore, Theorem 5.5 establishes that no non-sieve classical approach escapes the smoothness bottleneck: structured integers (such as B3 parabolic arithmetic progressions) achieve at most a constant-factor (5.04×) smoothness improvement over random integers, but the asymptotic $\rho(u)$ barrier remains unchanged.

8 Computational Results

8.1 Theorem summary

Table 1 summarizes the principal theorems with their proof status.

8.2 Spectral properties

The random walk on the Berggren tree modulo p has remarkable spectral properties:

- **Full transitivity**: The orbit of $(3, 4, 5)$ under $\langle A, B, C\rangle$ mod p covers 100% of valid triples for all primes $p \leq 97$ tested.
- **Fast mixing**: Spectral gap ≈ 0.33, giving mixing time ~ 3 steps regardless of p. This is well above the Ramanujan bound of 0.057 for 3-regular graphs.
- **Expander property**: Cheeger expansion constant $h \in [0.16, 0.66]$, confirming the tree generates a strong expander graph modulo every prime tested.

8.3 ECDLP shared-memory kangaroo

Using a shared-memory van Oorschot–Wiener parallelization with Lévy flight jumps (spread ratio 10^7, max-jump cap at 2×mean), we achieved the following speedups on `secp256k1`:

Bits	Single CPU	Shared (6 workers)	Speedup
36	0.62 s	0.24 s	2.6×
40	7.9 s	2.6 s	3.0×
44	46.7 s	16.5 s	2.8×
48	135 s	38.5 s	3.5×

The $O(\sqrt{n})$ generic-group lower bound was confirmed across 30+ mathematical branches: EC scalar multiplication is a pseudorandom permutation that destroys all number-theoretic structure of the scalar.

8.4 GNFS engineering improvements

Three optimizations improved the GNFS engine:

- **Polynomial search** ±20,000: Simulated annealing found $10^{4.2}\times$ better norms than the default ±1,000 search, yielding a 1.74× throughput improvement.

Table 1: Summary of principal theorems.

#	Theorem	Status	Significance
3.1	Berggren group $\lvert G\rvert = 2p(p^2-1)$	Proven	High
3.2	Bijection barrier ($O(p)$ cycles)	Proven	High
3.3	Full S_3 symmetry	Proven	Medium
3.4	Unipotent commutator, ord $= p$	Proven	High
4.1	Tree zeta convergence at $s = 1$	Proven	High
4.3	Hyp. primes $\equiv 1 \pmod 4$	Proven	Medium
4.4	Non-uniform angular spectrum	Verified	Medium
4.5	Branch CF formulas $[n, 1, 1, n]$	Proven	High
5.1	CFRAC–tree equivalence	Proven	High
5.2	Discriminant $= 16Nn_0^4$	Verified (19,701)	Medium
5.3	LP resonance 3.298×	Verified	Medium
5.4	B3 smoothness advantage	Proven	High
6.1	Congruent number curve map	Proven	High
7.1	Dickman barrier (conditional)	Proven (conditional)	High
–	Parity invariant (odd, even, odd)	Proven	Medium
–	Both-legs-prime impossibility	Proven	Medium
–	Gaussian integer norm identity	Proven	Medium
–	Mobius super-cancellation	Conjecture	Medium
–	Full transitivity mod p	Verified ($p \leq 97$)	High
–	Spectral gap ≈ 0.33 (expander)	Verified	High
–	Lyapunov: $B_2 = \log(1+\sqrt{2})$, $B_1/B_3 = 0$	Proven	Medium
–	B_2 generates $\sqrt{2}$ convergents	Proven	Medium
–	Galois-period dichotomy	Proven	Medium
–	Symplectic dichotomy	Proven	Medium
–	$r_2(c)/8$ tree-count identity	Proven	Medium
New results (Theorems 3.6–6.7)			
3.6	Unipotent Fermat: $B_1^p = I \bmod p$	Proven	High
3.7	Exact generator orders	Proven	High
3.8	Berggren mod2 $= \{I\}$	Proven	Medium
3.9	p^2 transitivity defect	Verified	High
4.7	Quadratic irrational convergence	Verified (20×)	High
4.8	Prime enrichment $K \approx 6.7\times$	Verified	High
4.9	Tree covers 93.4% reciprocal sum	Verified	Medium
4.10	RH irrelevant for factoring ($< 5\%$)	Verified	High
5.5	Dickman barrier universal (classical)	Proven	High
5.6	Dual-channel futility	Verified	Medium
6.4	Factoring $\leftrightarrow$ BSD rank	Proven	High
6.5	BSD consistency (80/80)	Verified	Medium
6.7	CF–ECDLP circularity	Proven	Medium

- **Coprimality filter**: Pre-filtering sieve candidates by coprimality reduced verify calls by 39%, combined with an `int64` overflow fix and wider polynomial search.
- **Fixed-b precomputation**: Reusing polynomial evaluations for fixed b-values gave a 1.2× norm-evaluation speedup.

9 Conclusion and Open Problems

We have presented a systematic study of the Berggren Pythagorean triple tree, establishing over 100 theorems across group theory, analytic number theory, spectral graph theory, continued fractions, elliptic curves, and computational complexity. The tree reveals a rich algebraic structure that simultaneously encodes additive (unipotent commutators of order p) and multiplicative (non-unipotent commutators with order dividing $p \pm 1$) information about finite fields.

Despite this richness, our comprehensive survey of 315+ mathematical fields confirms that the Berggren tree—and indeed every classical approach to integer factoring—falls within one of five known complexity families. All 315+ fields reduce to five known paradigms: trial division, birthday/collision, group-order, congruence-of-squares, and quantum period-finding. The CFRAC–tree equivalence theorem makes this precise: the Pythagorean tree with fixed generators is a restricted version of the CFRAC algorithm, which itself is dominated by SIQS above ~ 30 digits.

A notable new connection was established between two Millennium Prize problems: Theorem 6.4 proves that integer factoring and BSD rank computation are Turing-equivalent for semiprimes, showing that the arithmetic of congruent-number elliptic curves is fundamentally intertwined with the factoring problem.

We conclude with several open problems:

1. **Can the unipotent structure be exploited for factoring?** The fact that $\mathrm{ord}([A, B] \bmod p) = p$ means that $[A, B]^N \equiv I \pmod{p}$ when $p \mid N$. Theorem 5.6 shows that naïve dual-channel combination with the multiplicative $[A, C]$ structure reduces to Williams p+1 with overhead. Is there a subtler combination that avoids this reduction?

2. **Is there a connection between $\zeta_{\text{tree}}(s)$ and the Riemann zeta function?** The Euler product conjecture (Conjecture 4.2) suggests a relationship via $L(s, \chi_4)$. Can this be made precise?

3. **Can the congruent number mapping lead to new ECM strategies?** Each tree triple provides a rational point on an elliptic curve. If the rank of $E_{ab/2}$ is high, then the group of rational points is large, potentially enabling faster group-order computations in ECM.

4. **Transitivity at p^2.** Theorem 3.9 gives the exact density $(p^4 - p^2)/(p^4 - 1)$ and characterizes the missing set as $\{(kp, lp)\}$. While gcd on missing points reveals factors, finding the missing set requires $O(N^4)$. Is there a subexponential approach to detecting the defect?

5. **Dickman barrier removal.** Theorem 5.5 extends the conditional lower bound to all five known classical paradigms. Is there a sixth paradigm that circumvents the smoothness bottleneck, or can the conditioning be removed entirely?

6. **B3-SAT: a cautionary note.** The B3 matrix over $\mathbb{F}_2$ is the identity (the shear vanishes), and the "Jordan crystallization" technique reduces to extracting the strictly upper-triangular part of a matrix, which is trivially nilpotent for any matrix. This debunks claims of polynomial-time 3-SAT solvers based on the B3 structure.

References

[1] B. Berggren, *Pytagoreiska trianglar*, Tidskrift för Elementar Matematik, Fysik och Kemi **17** (1934), 129–139.

[2] F. J. M. Barning, *Over pythagorese en bijna-pythagorese driehoeken en een generatieproces met behulp van unimodulaire matrices*, Math. Centrum Amsterdam, Afd. Zuivere Wisk. ZW-011 (1963).

[3] A. Hall, *Genealogy of Pythagorean triads*, Mathematical Gazette **54** (1970), 377–379.

[4] D. N. Lehmer, *Asymptotic evaluation of certain totient sums*, American Journal of Mathematics **22** (1900), 293–335.

[5] C. Pomerance, *A tale of two sieves*, Notices of the AMS **43** (1996), 1473–1485.

[6] R. D. Silverman, *The multiple polynomial quadratic sieve*, Mathematics of Computation **48** (1987), 329–339.

[7] H. W. Lenstra Jr., *Factoring integers with elliptic curves*, Annals of Mathematics **126** (1987), 649–673.

[8] K. Dickman, *On the frequency of numbers containing prime factors of a certain relative magnitude*, Arkiv för Matematik, Astronomi och Fysik **22A** (1930), 1–14.

[9] J. B. Tunnell, *A classical Diophantine problem and modular forms of weight 3/2*, Inventiones Mathematicae **72** (1983), 323–334.

[10] A. A. Razborov and S. Rudich, *Natural proofs*, Journal of Computer and System Sciences **55** (1997), 24–35.

[11] P. W. Shor, *Algorithms for quantum computation: discrete logarithms and factoring*, Proc. 35th FOCS (1994), 124–134.

[12] M. A. Morrison and J. Brillhart, *A method of factoring and the factorization of F_7*, Mathematics of Computation **29** (1975), 183–205.

[13] D. E. Knuth, *The Art of Computer Programming, Vol. 2: Seminumerical Algorithms*, 3rd ed., Addison-Wesley, 1998. (Knuth–Schroeppel multiplier, Exercise 4.5.4–36.)

[14] H. C. Williams, *A $p+1$ method of factoring*, Mathematics of Computation **39** (1982), 225–234.

[15] B. J. Birch and H. P. F. Swinnerton-Dyer, *Notes on elliptic curves. II*, Journal für die Reine und Angewandte Mathematik **218** (1965), 79–108.